

TensorFlow Guide

Learning Outcomes

By the end of this ride, you'll be able to:

- Spin up a dedicated TensorFlow workspace like a boss—no stray files messing with your vibe.
- Compile and install Python 3.11 from source, flexing those optimization muscles.
- Create and toggle virtual environments, isolating your experiment like a mad scientist.
- Fetch and unpack a pretrained SSD-MobileNet model + COCO label map, so you're not reinventing the wheel.
- Wire up Flask, OpenCV, NumPy, and TensorFlow into a live webcam-powered object detector.

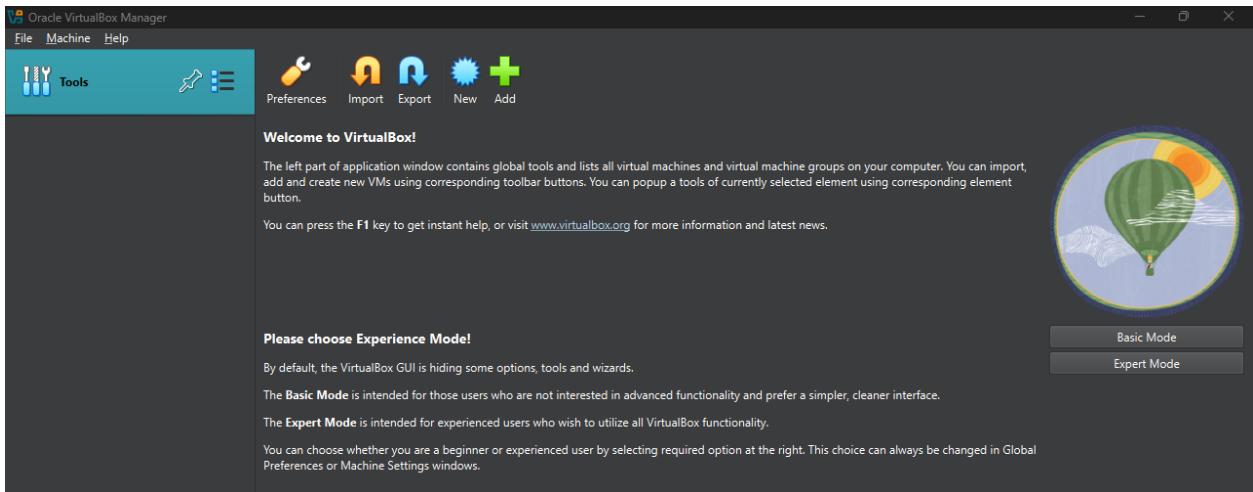
Prerequisites

- An Ubuntu VM goto <https://lchs.mm.fcix.net/ubuntu-releases/25.04/ubuntu-25.04-desktop-amd64.iso>
- A vm with access to your webcam (ideally a usb webcam or built in)
- At least 200mb of free space
- VM specs (ideally)
 - **CPU (x86)**
 - **Min:** 4 vCPUs (2.0 GHz+)
 - **Sweet spot:** 8 vCPUs or more
 - **RAM**
 - **Min:** 8 GB
 - **Reco:** 16 GB+
 - **Other**
 - Git and Curl Installed
 - **Display**
 - 320x320 minimum

Setting up the Ubuntu VM

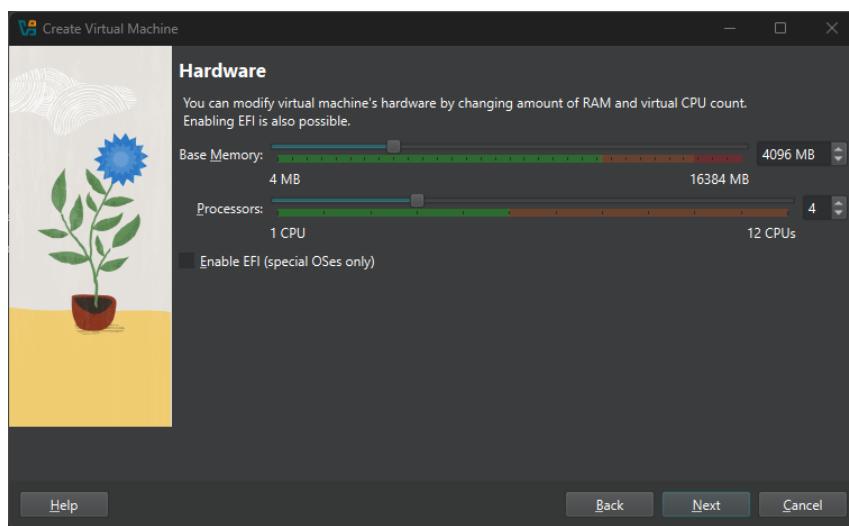
1. Go to <https://lchs.mm.fcix.net/ubuntu-releases/25.04/ubuntu-25.04-desktop-amd64.iso> to download the iso file
2. Go to <https://www.virtualbox.org/wiki/Downloads>, Install and setup Virtualbox.

You should get a result like this:

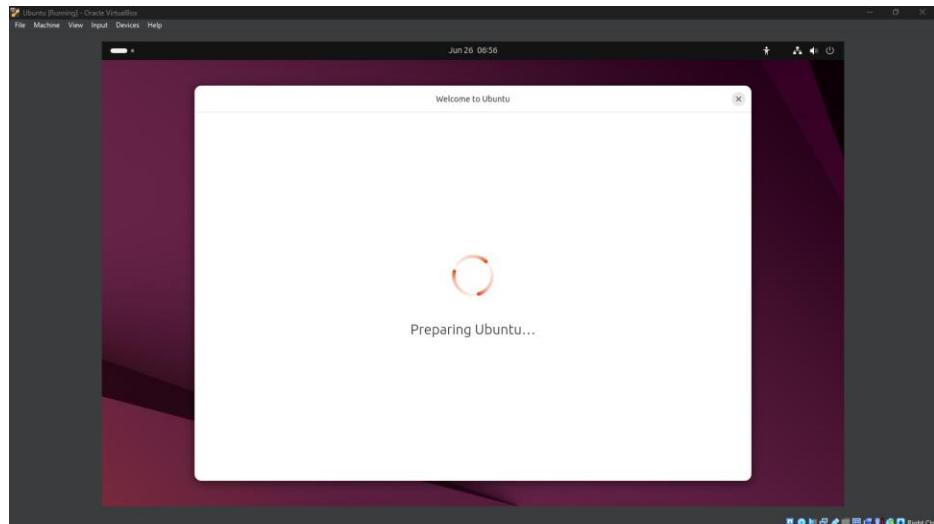


3. Following along, press the new button and follow instructions to setup the vm

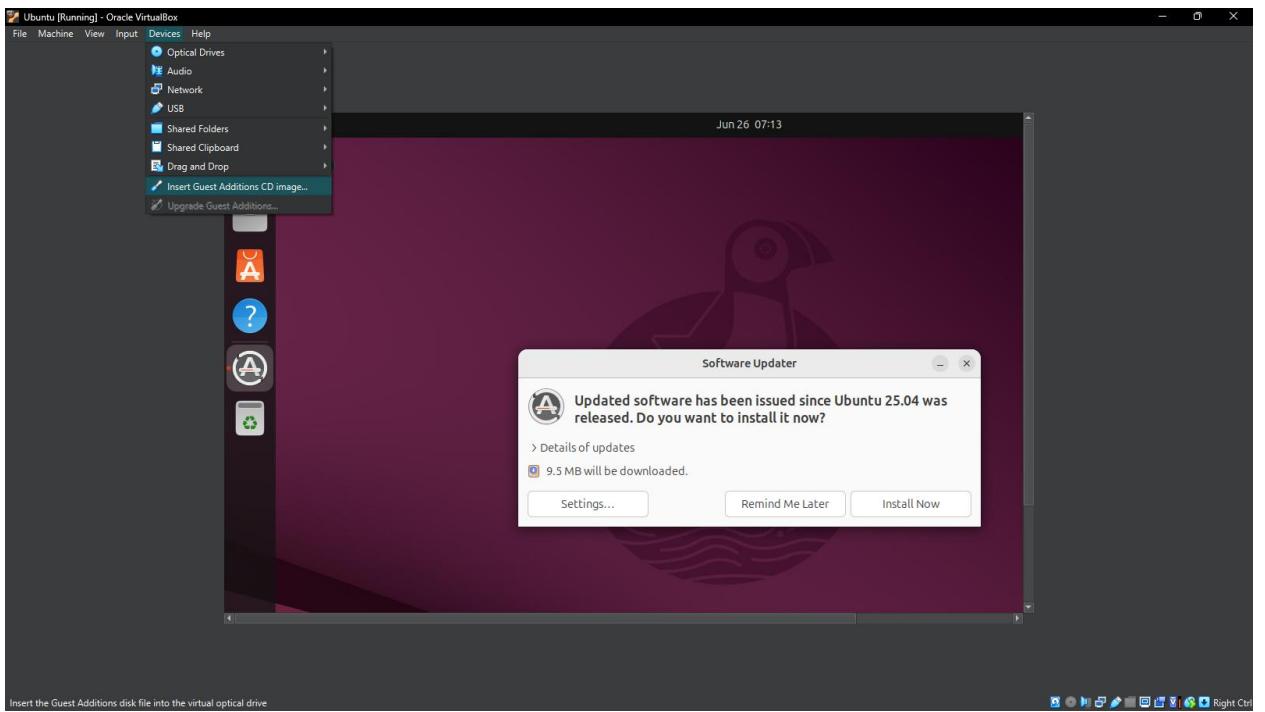
Note: allow about 4 GB of RAM and 4 CPUs to run



- VM should start running. Press Try and install Ubuntu in grub and follow along the unintended GUI installation

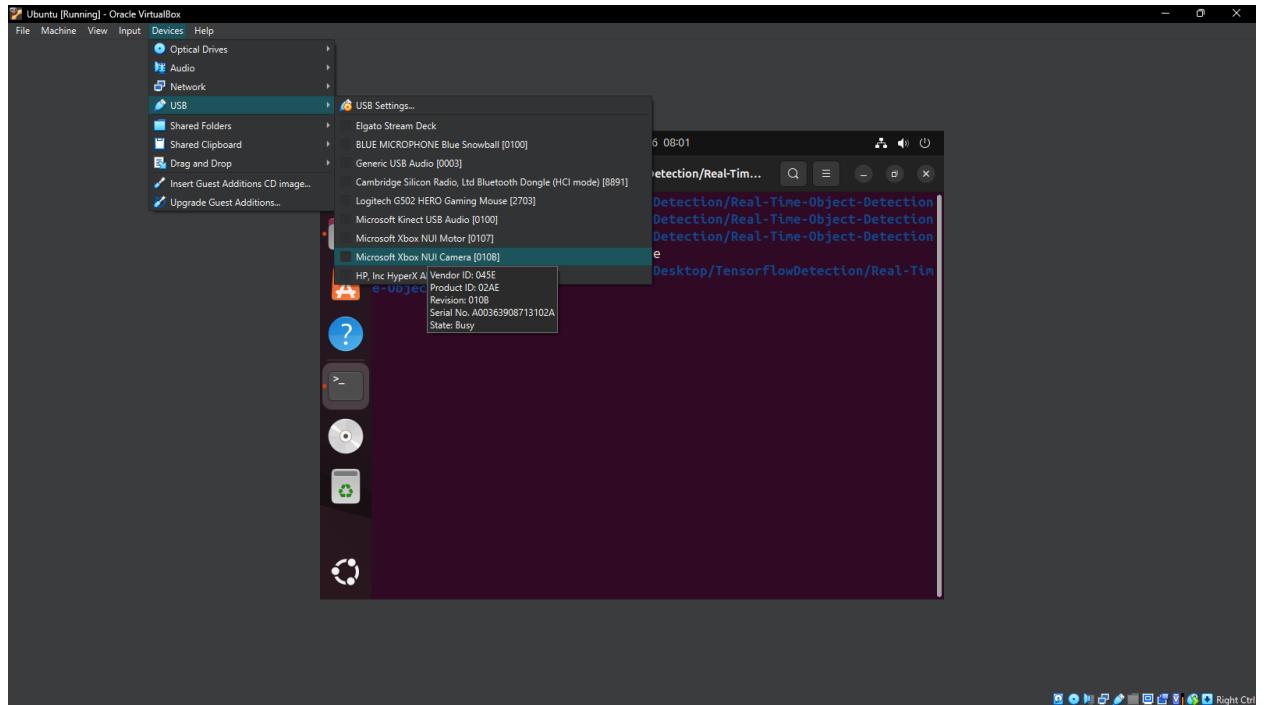


- Optional: Once the vm reboots, feel free to install the Guest Additions for a smoother experience (& shared clipboard)



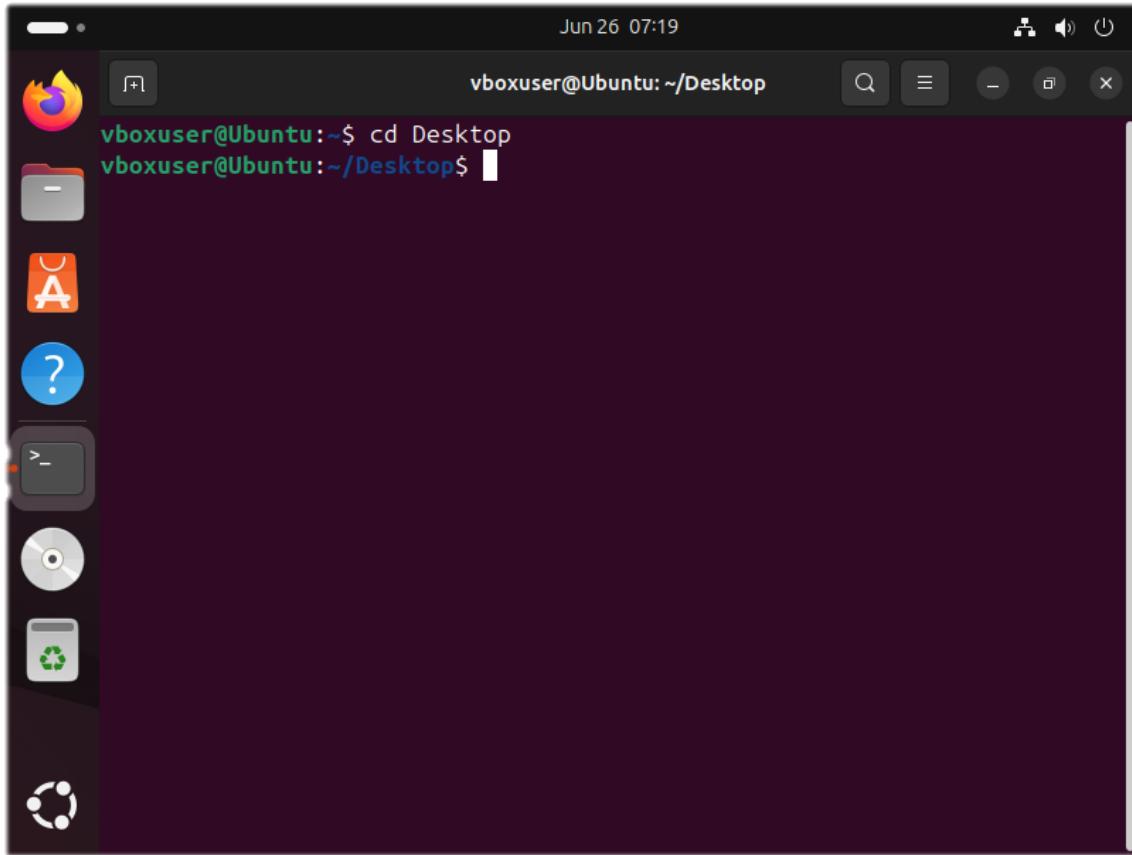
Enable shared clipboard at [Devices > Shared Clipboard > Bidirectional](#)

6. Now enable usb webcam passthrough.



Setting up Folders

Now on the desktop, open up terminal (**ctrl+alt+t**), and enter **cd Desktop** to the terminal to get to the desktop.



Now run the following to setup git and curl, for later.

```
sudo apt install git  
sudo apt install curl
```

Now, to setup the file system for this project.

```
mkdir TensorflowDetection  
cd TensorflowDetection  
git clone https://github.com/Yonas650/Real-Time-Object-Detection.git  
cd Real-Time-Object-Detection
```

Why we do this:

- **mkdir TensorflowDetection:** sets up your own lil' sandbox—keeps your files from minglin' with your desktop.
- **cd TensorflowDetection:** jump into that sandbox.

- **git clone ...**: grabs the entire Real-Time-Object-Detection codebase so you can hack on it.
- **cd Real-Time-Object-Detection**: enters the project folder, where all the magic lives.

Final Result:

```

vboxuser@Ubuntu: ~/Desktop/TensorflowDetection/Real-Tim...
vboxuser@Ubuntu: ~/Desktop$ cd Desktop
vboxuser@Ubuntu: ~/Desktop$ mkdir TensorflowDetection
cd TensorflowDetection
git clone https://github.com/Yonas650/Real-Time-Object-Detection.git
cd Real-Time-Object-Detection
Cloning into 'Real-Time-Object-Detection'...
remote: Enumerating objects: 73, done.
remote: Counting objects: 100% (73/73), done.
remote: Compressing objects: 100% (50/50), done.
remote: Total 73 (delta 39), reused 50 (delta 20), pack-reused 0 (from 0 )
Receiving objects: 100% (73/73), 15.60 Kib | 7.80 MiB/s, done.
Resolving deltas: 100% (39/39), done.
vboxuser@Ubuntu: ~/Desktop/TensorflowDetection/Real-Time-Object-Detection
$
```

Setup Update & Dependencies

Now its time to install more prerequisites.

```

sudo apt update

sudo apt install -y wget build-essential zlib1g-dev libncurses5-dev \
libgdbm-dev libnss3-dev libssl-dev libreadline-dev libffi-dev \
libsqlite3-dev \
libbz2-dev liblzma-dev
```

What's poppin' here:

- **sudo apt update**: syncs your package list with the interwebs—keeps you fresh.
- **sudo apt install -y ...**: grabs core tools and libraries you need to compile Python from scratch.
 - *build-essential*: GCC, make, all that compile jazz.

- *zlib, ncurses, openssl, sqlite...*: these let Python support compression, terminal UIs, encryption, and that sweet built-in database.

Download & Build Python 3.11.9

TensorFlow requires an older version of Python over the new one, because of this , we ill be using an older version.

```
cd /tmp
wget https://www.python.org/ftp/python/3.11.9/Python-3.11.9.tgz
tar -xf Python-3.11.9.tgz
cd Python-3.11.9
./configure --enable-optimizations
make -j$(nproc)
sudo make altinstall
cd ~/Desktop/TensorflowDetection/Real-Time-Object-Detection
```

Deets:

- **cd /tmp**: temp folder so your home dir stays clean.
- **wget ...Python-3.11.9.tgz**: download the source tarball of the latest Python you wanna rock.
- **tar -xf ...**: unpacks that tarball.
- **./configure --enable-optimizations**: sets flags to squeeze out performance.
- **make -j\$(nproc)**: build using all your CPU cores—parallel hustle.
- **sudo make altinstall**: installs as python3.11 so you don't stomp your system Python.

Now check your python version with `python3.11 --version`

```
vboxuser@Ubuntu:~/Desktop/TensorflowDetection/Real-Time-Object-Detection
$ python3.11 --version
Python 3.11.9
```

Setup & Activate Virtual Environment

```
python3.11 -m venv tensorflow-venv  
source tensorflow-venv/bin/activate  
pip install --upgrade pip
```

Why VENV?

- **Isolates dependencies:** no clashes with global packages.
 - **pip install —upgrade pip:** keeps pip on fleek.

Make sure you have (tensorflow-venv) in your terminal, this means ur in the virtual environment.

```
(tensorflow-venv) vboxuser@Ubuntu:~/Desktop/TensorflowDetection/Real-Time-Object-Detection$
```

Install Python Libraries

What you get:

- **TensorFlow**: the neural-net engine.
 - **OpenCV**: vision ops—reading frames & drawing boxes.
 - **NumPy**: tensor math under the hood.
 - **Flask**: lightweight web server to serve your camera stream.

Note: This might take a while....

```
Downloading tensorflow-2.19.0-cp311-cp311-manylinux_2_17_x86_64.manylinu  
x2014_x86_64.whl (644.9 MB) 240.1/644.9 MB 5.6 MB/s eta 0:01:13
```

Installing Model and Label Map

```
wget  
http://download.tensorflow.org/models/object_detection/tf2/20200711/ssd  
_mobilenet_v2_fpnlite_320x320_coco17_tpu-8.tar.gz -O  
ssd_mobilenet.tar.gz  
  
tar -xzvf ssd_mobilenet.tar.gz  
  
curl -o mscoco_complete_label_map.pbtxt  
https://raw.githubusercontent.com/tensorflow/models/master/research/obj  
ect_detection/data/mscoco_complete_label_map.pbtxtcd  
  
cd ~/Desktop/TensorflowDetection/Real-Time-Object-Detection
```

Note: do each command one by one

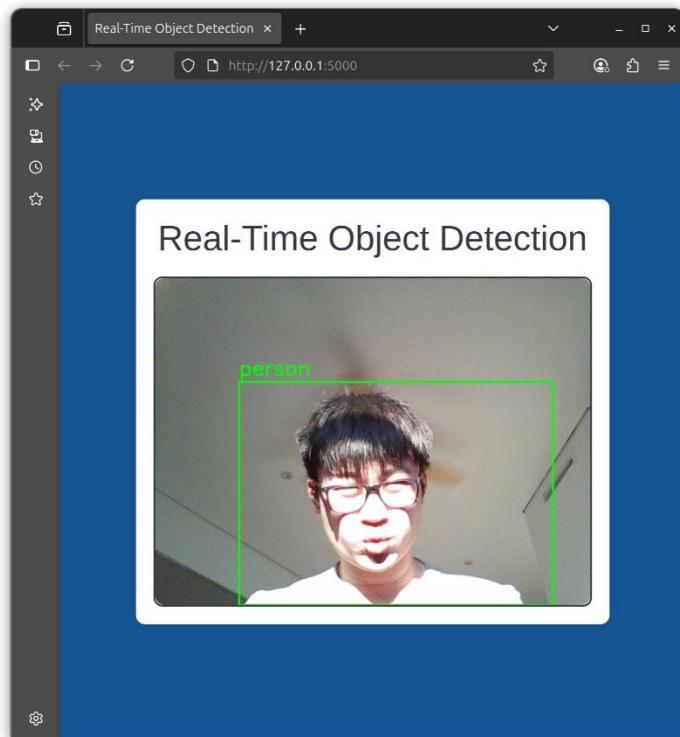
Breakdown:

- **wget ...ssd_mobilenet...:** pulls the SSD-MobileNet model pre-trained on COCO so you don't train from scratch.
- **tar -xzvf:** unpacks the model files (checkpoints, pipeline config).
- **curl -o ...label_map.pbtxt:** downloads the label map—maps integer IDs to human labels (e.g., “person”, “dog”).

Run and pray 🙏

```
python3.11 model.py  
firefox http://127.0.0.1:5000 & # in a new terminal tab
```

Note: if you see a window like this:



Congrats! Youve Reached the end!